MemForge: A Deterministic Simulator for Operating-System Memory Management

An OS Mini Project submitted for the partial fulfillment of
the requirements for the degree

of

BACHELOR OF TECHNOLOGY

in

COMPUTER SCIENCE & ENGINEERING

By

Adrija Das (12024002001001)

Meghna Dhar (12024002001045)

Arpita Goswami (12024002001021)

Ankita Dhara (12024002001263)

Pritam Majumder (12024002001303)

Anusuya Sarkar (12024002001179)

Jayasree Saha (12024002001289)

Aloke Dipto Banerjee (12024002001006)

Sandipan Gain (12024002001232)

Arya Gupta (12024002001189)

TEAM NO. 5

Department of Computer Science and Engineering

Institute of Engineering and Management

(School of University of Engineering and Management, Kolkata)

Newtown Sector

2026

DECLARATION

We hereby declare that this project is the product of our own independent work and has not been submitted, in whole or in part, to any other university or institution for the award of any degree, diploma, or other qualification. All sources of information have been duly acknowledged, and any work carried out in collaboration with others has been clearly indicated in the text.

Signed by the B.Tech Students: -

Sl. No.	**Name**	**Class Roll No.**	**Signature**
1	Adrija Das	51	
2	Meghna Dhar	52	
3	Arpita Goswami	53	
4	Ankita Dhara	54	
5	Pritam Majumder	55	
6	Anusuya Sarkar	56	
7	Jayasree Saha	57	
8	Aloke Dipto Banerjee	58	
9	Sandipan Gain	59	
10	Arya Gupta	60	

Date: -

ACKNOWLEDGEMENT

We would like to take this opportunity to express our heartfelt gratitude to everyone who inspired, supported, and contributed — directly or indirectly — to the successful completion of this OS mini project.

First and foremost, we extend our deepest thanks to our families for their unwavering support and encouragement throughout our academic journey. We are equally grateful to our project guide, [Guide Name], for his/her unwavering support, guidance, and encouragement throughout this project. His/her knowledge and guidance have been invaluable in shaping our work and in preparing this report. We would also like to express our sincere thanks to the faculty members and fellow students of the Department of Computer Science and Engineering, University of Engineering and Management, Kolkata, for their helpful suggestions and support, and for allowing us to call this institution home.

Finally, we bow in gratitude to the Almighty for the strength, clarity, and perseverance to endure and overcome the challenges along this journey. This project has been a rewarding learning experience, and we feel deeply blessed to have had the support and grace needed to complete it.

Team No. 5
University of Engineering and Management, Kolkata

TABLE OF CONTENTS

S. No.	Section Topic	Page No.
1	Introduction	5
2	Objectives, Scope and Limitations	5
3	Operating System Memory Management Concepts	5
4	System Architecture and Design	7
5	Module Description and Dependencies	9
6	Memory Model and Allocation State	9
7	First-Fit, Best-Fit and Worst-Fit Algorithms	10
8	Fragmentation and Compaction	11
9	Memory Allocation Examples and Results	11
10	Paging Model and Address Translation	12
11	FIFO, LRU and Optimal Replacement	13
12	Page Replacement Worked Examples	13
13	Paging Results and Comparison	14
14	Workloads and Performance Metrics	15
15	Benchmarking and Performance Analysis	15
16	Testing Strategy and Correctness	16
17	Test Cases and Test Results	17
18	User Interface and Workflow	17
19	Execution Examples and Output	18
20	Repository Structure and Development Conventions	18
21	Build System and Continuous Integration	18
22	Setup Instructions and Documentation	19
23	Conclusion and Future Scope	19
24	References	19
25	Appendix	20

MemForge is a deterministic simulator of an operating-system memory manager, written in C17 and also compiled to WebAssembly. It covers contiguous allocation (First-, Best- and Worst-Fit, splitting, coalescing, sliding compaction) and demand paging (single-level page tables, address translation, FIFO, LRU and OPT replacement). Policies only decide and engines only mutate, and every change is published as an event, so identical inputs and seeds always produce byte-identical traces. On 30 seeded workloads at 90% utilisation, Best-Fit reaches a 91.4 ± 0.4% success rate against 90.0 ± 0.6% for First-Fit and 87.2 ± 0.7% for Worst-Fit, while First-Fit scans the fewest holes (3.21 vs. 5.24 per request). On 2,000-reference workloads, faults follow OPT < LRU < FIFO at every frame count (103.2, 138.9 and 253.1 at 8 frames), and a frame sweep reproduces Belady's anomaly (9 → 10 faults). All eight acceptance criteria and 17 test suites pass under AddressSanitizer, UBSan and Valgrind.

***Index Terms*—memory management, deterministic simulation, contiguous allocation, fragmentation, demand paging, page replacement, Belady's anomaly.**

I. Introduction

Every running program (a *process*) needs space in the computer's main memory (RAM) to hold its code and data. RAM is limited and shared by many processes, so the operating system must decide *who gets which part of memory*, give it back when a process finishes, and stop processes from touching each other's memory. This job is called *memory management*.

Memory management decides where each process lives in physical memory and how its addresses are translated. Textbooks present contiguous allocation and demand paging with small hand-worked examples, but the trade-offs between policies only become visible when they run on the same workload, step by step, at scale.

MemForge closes that gap. It is a modular C17 simulator with a command-line tool for scripted experiments and a browser laboratory that runs the *same* engine compiled to WebAssembly. Its three design pillars are summarised in Fig. 1; the rest of this paper follows the order in which the system was built: concepts, architecture, the contiguous subsystem, the paging subsystem, experiments, testing, interfaces and engineering.

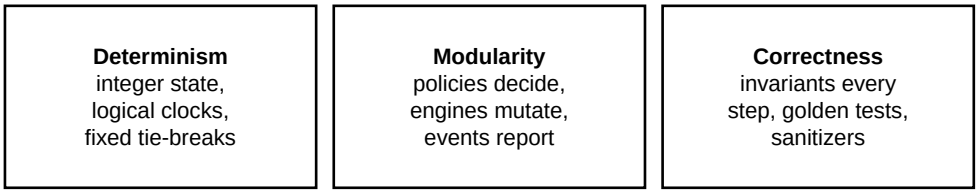


Fig. 1. The three design pillars of MemForge.

II. Objectives, Scope and Limitations

A. Objectives

1. Implement First-, Best- and Worst-Fit with splitting, coalescing, fragmentation metrics and compaction.
2. Simulate translation, page faults and FIFO, LRU and OPT replacement.
3. Compare policies fairly on identical, reproducible workloads.
4. Verify correctness with invariants, automated tests and memory-safety tools.
5. Offer a CLI and an interactive WebAssembly lab.

B. Scope

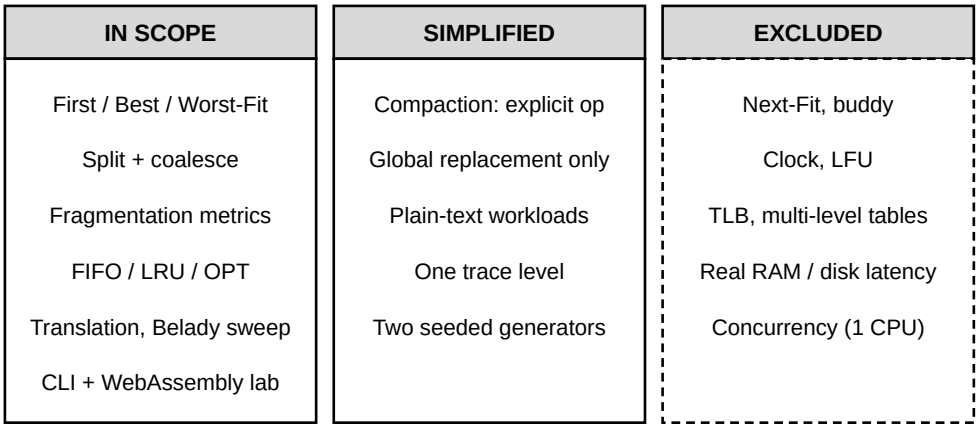


Fig. 2. Scope boundaries: what is implemented, simplified and excluded.

C. Limitations

Memory is simulated, not real RAM; one CPU applies one operation at a time; disk I/O is counted but its latency is not modelled; page tables are single-level with no TLB; Next-Fit, Clock and other policies are out; compaction is never automatic.

III. Operating System Memory Management Concepts

From the ground up. Think of memory as one long row of numbered boxes (addresses 0, 1, 2, …). A program asks for, say, 212 boxes; the OS must find 212 *free boxes in a row*, mark them as used, and later mark them free again. Two broad strategies exist: give each request one unbroken run of boxes (*contiguous allocation*), or cut both programs and

memory into equal-sized pieces that can be placed anywhere (*paging*). MemForge simulates both.

The MMU maps each *logical* address produced by the CPU to a *physical* address and traps invalid accesses. Under contiguous allocation each request receives one run of memory taken from a free *hole*: First-Fit takes the first hole that fits, Best-Fit the smallest and Worst-Fit the largest. *External* fragmentation is free memory split into holes too small to use; *internal* fragmentation is unused space inside a block (zero in MemForge, since granted = requested). Compaction slides blocks together to merge the holes (Fig. 3).

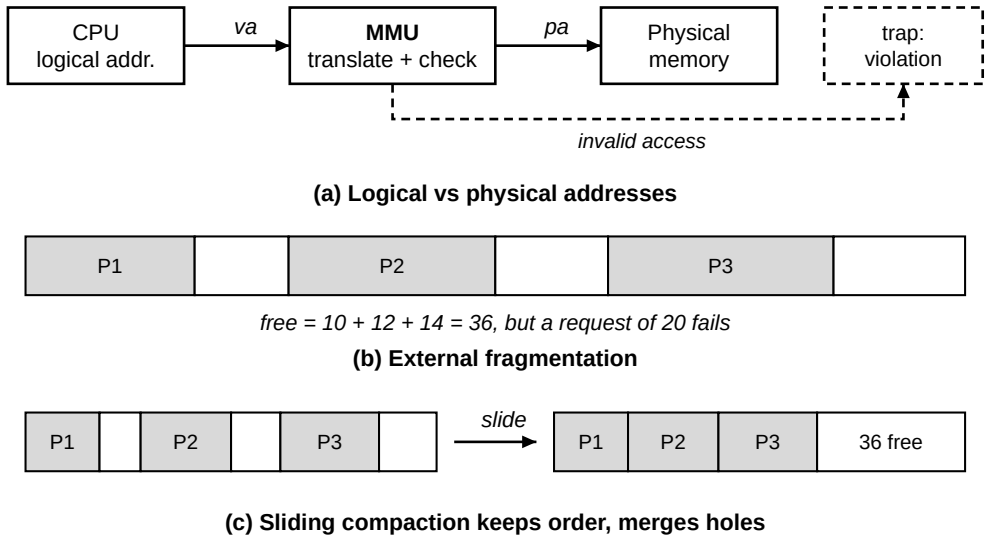


Fig. 3. Address translation, external fragmentation and compaction.

Why paging? With contiguous allocation, free space ends up scattered in small gaps. Paging avoids this: every piece is the same size, so any free slot fits any piece. The cost is a lookup table per program that records where each piece currently lives.

Paging splits logical memory into fixed-size *pages* and physical memory into equal *frames*; a page table maps one to the other (Fig. 4). Under demand paging a page is loaded only on its first use; a reference to a non-resident page is a *page fault*. When no frame is free, a replacement policy picks a victim: FIFO the oldest page, LRU the least recently used, OPT the page used furthest in the future. Belady's anomaly—more frames, more faults—can occur with FIFO but not with stack algorithms such as LRU and OPT (Fig. 5).

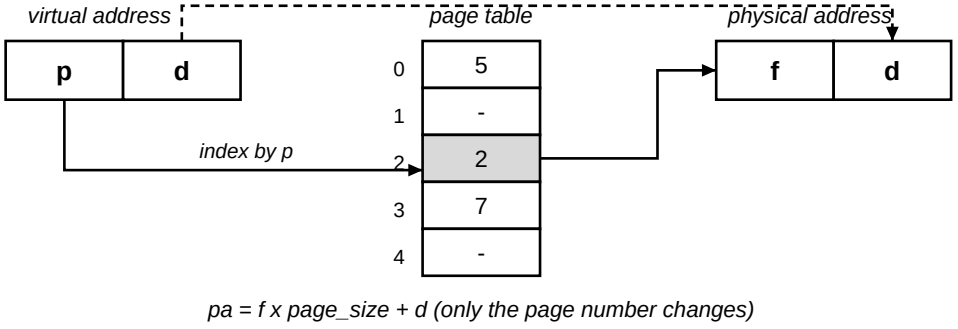


Fig. 4. Paging: the page number is replaced by a frame number.

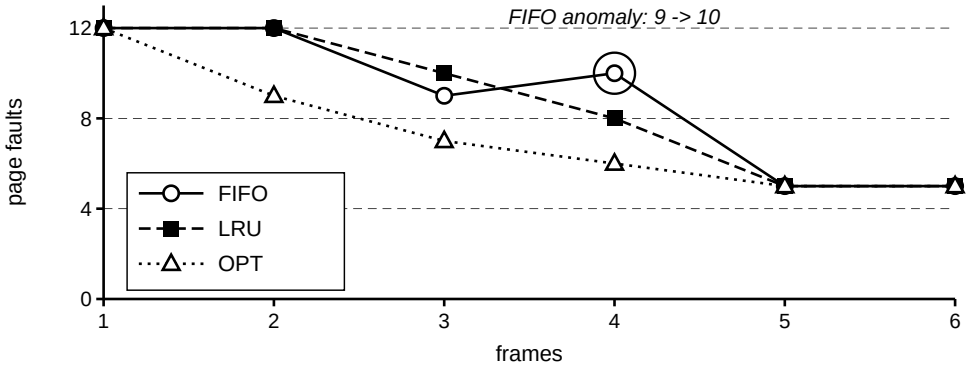


Fig. 5. Faults vs. frames on the string 1 2 3 4 1 2 5 1 2 3 4 5.

IV. System Architecture and Design

A. Layered Architecture

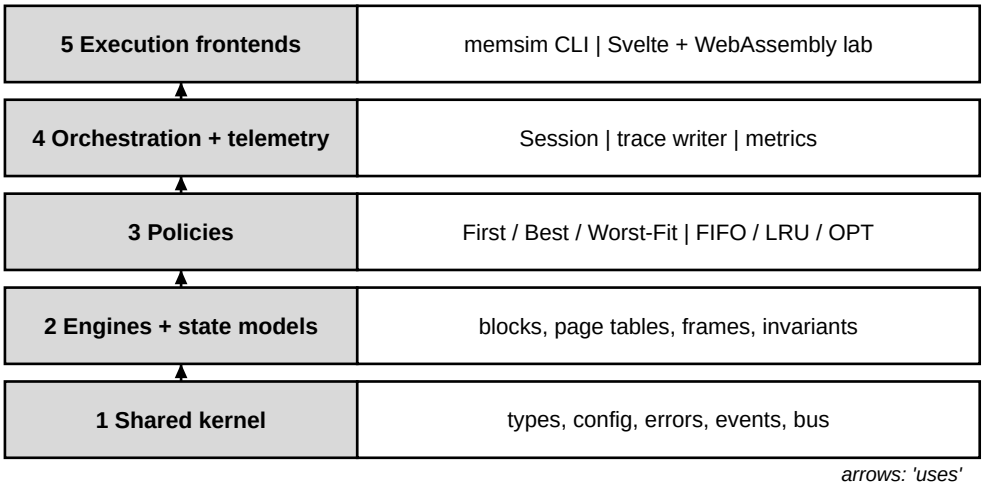


Fig. 6. Five layers; each layer uses only the ones below it.

B. Deterministic Core and Session Facade

The core performs no I/O and reads no wall clock; it uses logical clocks (*step*, *t*, *seq*), integer arithmetic and fixed tie-breaks (lowest address, lowest frame). Every client goes through the `Session` facade, which applies one operation, checks invariants and publishes the step's events in `seq` order (Fig. 7).

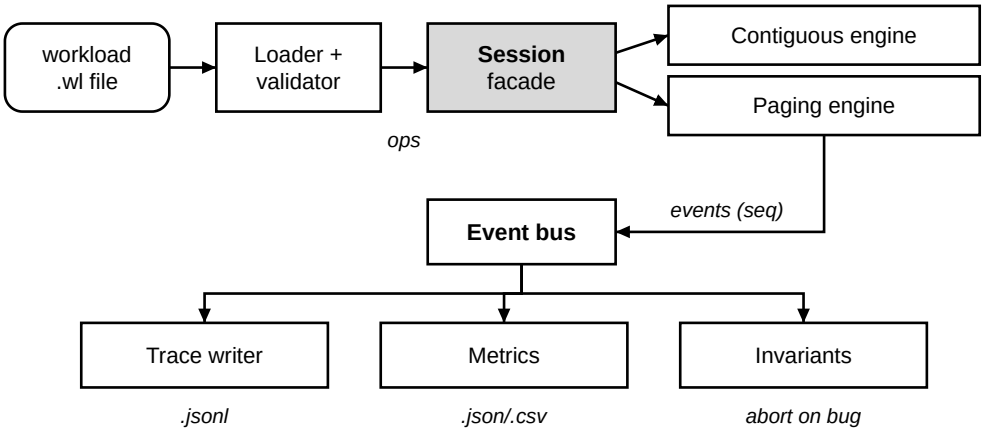


Fig. 7. End-to-end data flow of one run.

C. Events, Traces and WebAssembly Bridge

Each step yields events (23 types, Appendix) that the trace writer stores as JSON Lines with a snapshot after every step. The same C modules are compiled with Clang for the WASI target and loaded by the web app, so browser and CLI traces are byte-identical.

TABLE I
ARCHITECTURE DECISION LOG

Decision	Reason
C17	systems language, no runtime
Address-ordered block array	simple, deterministic
Policies separate from engines	testable, cannot corrupt state
Event log as contract	one source for trace, metrics, tests
Integer-only state	identical output on every machine
Global frame pool	matches textbook examples
Plain-text workloads	no JSON parser in C

D. Component Workflows

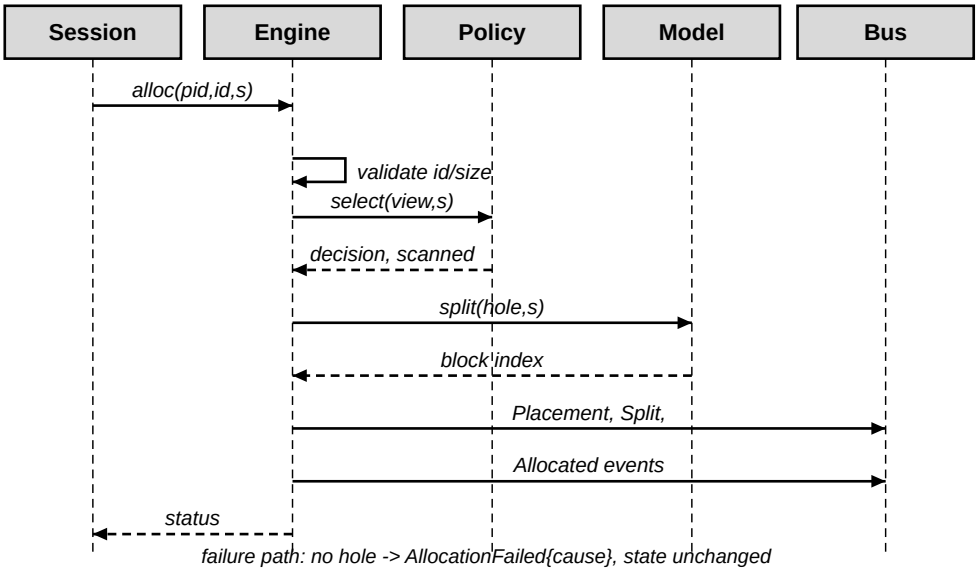


Fig. 8. Allocation sequence in the contiguous subsystem.

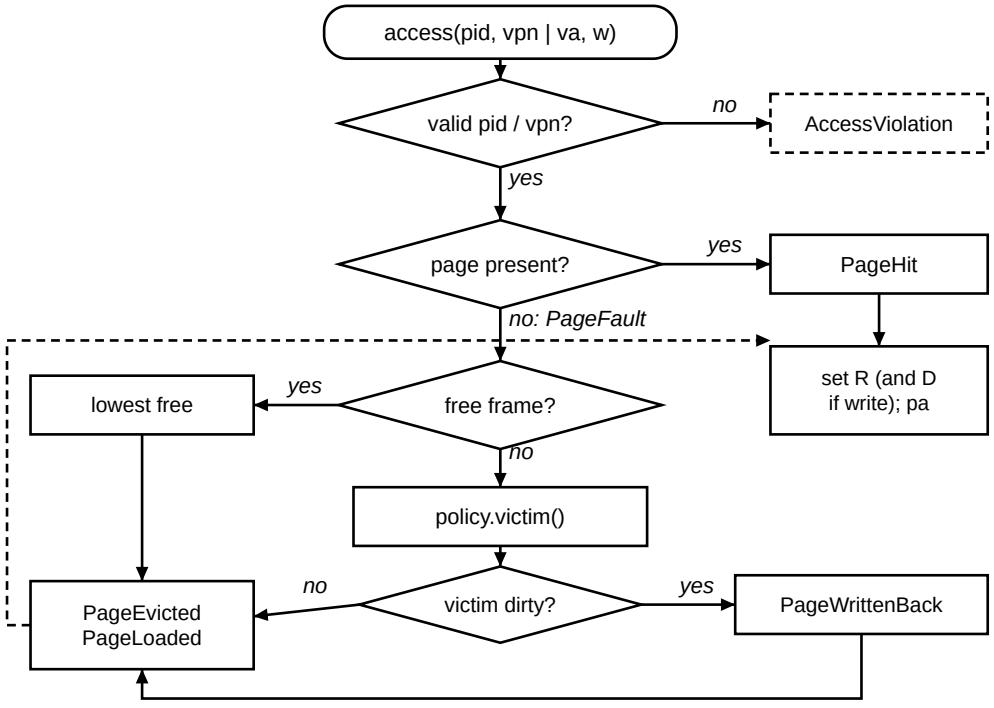


Fig. 9. Paged memory access with fault handling and replacement.

V. Module Description and Dependencies

TABLE II
Module catalogue

Module	Responsibility
common	types, config, errors, events, bus, registry
contiguous	blocks, engine, FF/BF/WF, fragmentation, compaction
paging	page/frame tables, MMU, FIFO/LRU/OPT, Belady sweep
trace	JSONL writer and snapshots
workloads	loader, validator, seeded generators
metrics	every reported figure, from events
session	facade: apply, check, publish
bench	runner over workload × policy × seed
cli	memsim commands, SHA-256

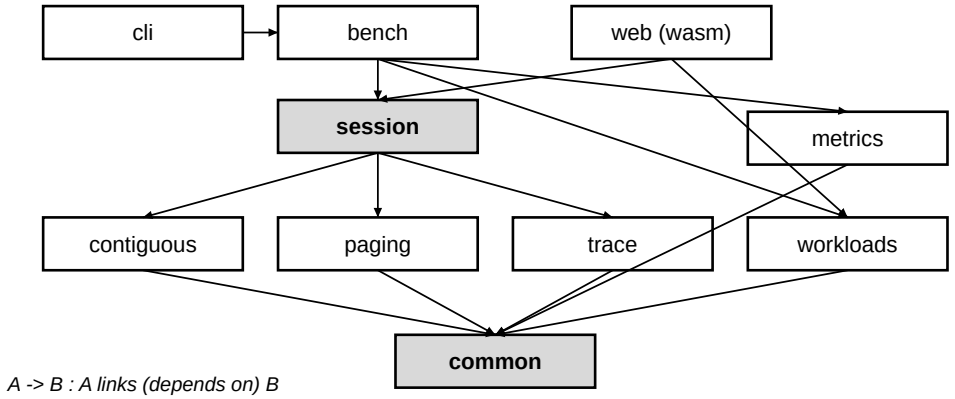
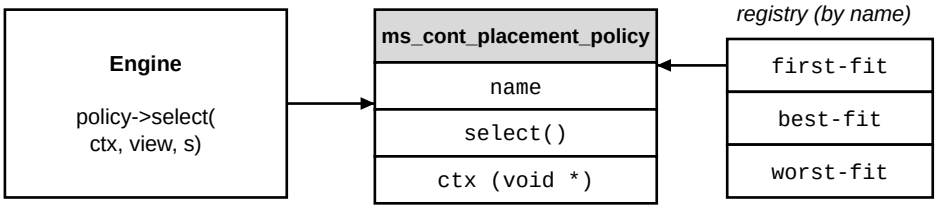


Fig. 10. Module dependency graph (acyclic, checked by check_layers.sh).

Each module exports only what others need; policies cannot modify engine data, and the core may not include `stdio.h`, `time.h` or system headers. Policies are function-pointer tables registered by name (Fig. 11), so a new policy is one file and one registry line.



policy reads a const MemoryView, returns a decision;
only the engine mutates state. FIFO/LRU/OPT use the same pattern.

Fig. 11. Policy vtable and registry.

VI. Memory Model and Allocation State

In plain terms, MemForge keeps a list of memory pieces in address order; each piece is either *used* (owned by a process) or *free* (a hole). Allocation cuts a hole in two; freeing joins a piece with free neighbours so holes never sit side by side.

Memory [0, M) is an address-ordered array of `ms_block` records that tile it exactly. Workloads refer to blocks by stable allocation ids, never by address, because compaction moves blocks; each id passes through the lifecycle in Fig. 12, which rules out duplicate ids and double frees.

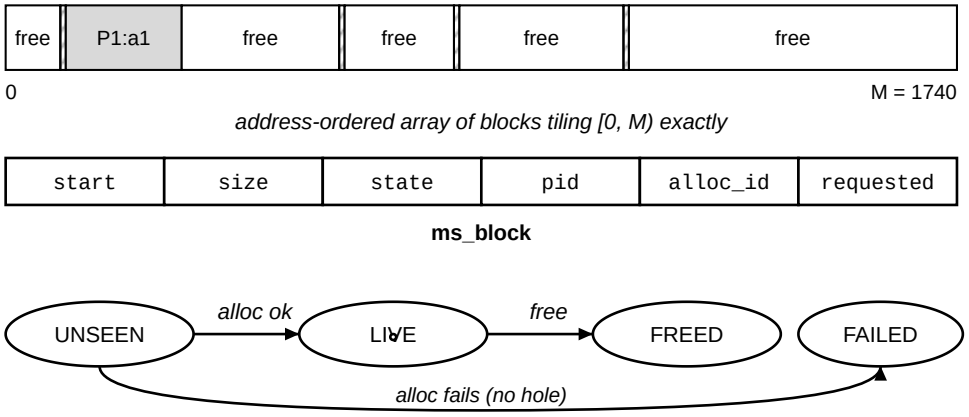


Fig. 12. Block array, block record and allocation-id lifecycle.

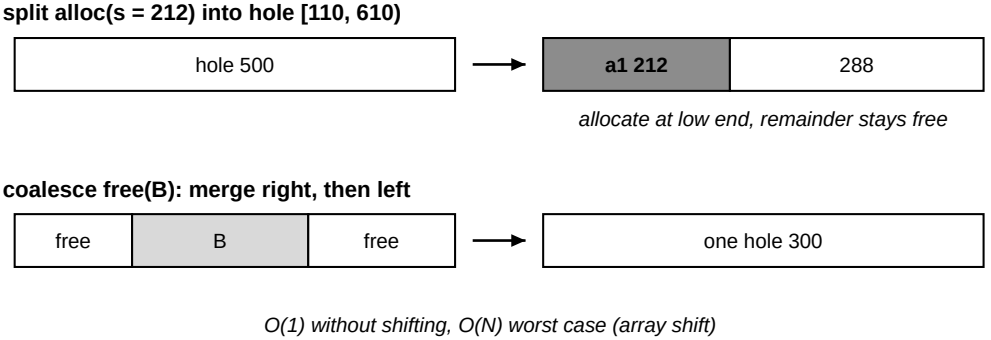


Fig. 13. Splitting a hole on allocation and coalescing on free.

TABLE III
MODEL OPERATIONS

Operation	Time	Space
split	O(1), O(N) worst (shift)	O(1)
free + coalesce	O(1), O(N) worst	O(1)

VII. FIRST-FIT, BEST-FIT AND WORST-FIT ALGORITHMS

When several holes are big enough, which one should be used? Picture parking a car: take the *first* space that fits (First-Fit), the *tightest* space that fits (Best-Fit), or the *biggest* space (Worst-Fit). Each choice affects how usable the leftover gaps will be later.

A placement policy is a pure selector: given the holes in address order and a size s, it returns whether a hole was found, which one, and how many holes it *scanned* (the search-cost metric). Algorithm 1 gives all three; they differ only in the selection test.

Algorithm 1: Placement selection (FF / BF / WF)

```
1: if view = NULL or s ≤ 0 then return (found = false, scanned = 0)
2: best ← none
3: for i ← 0 to h − 1 do
4:   if holes[i].size < s then continue
5:   FF: return (true, i, scanned = i + 1)
6:   BF: if best = none or size < holes[best].size then best ← i
7:   WF: if best = none or size > holes[best].size then best ← i
8: end for
9: return (best ≠ none, best, scanned = h)
```

Strict comparisons keep the lowest address on ties. All three run in O(h) time and O(1) space. Fig. 14 shows them on the textbook five-hole memory.

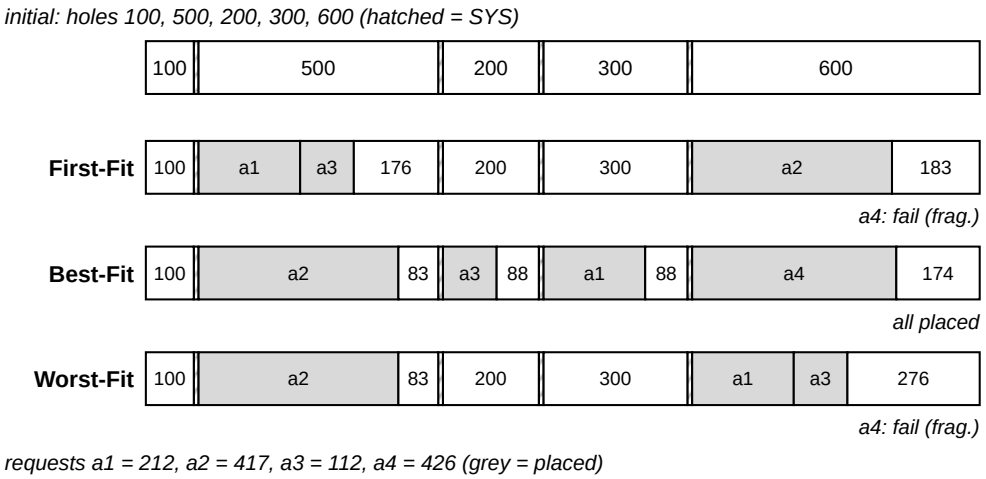


Fig. 14. Placements on the Silberschatz five-hole memory (AC-1).

Excluded: Next-Fit (a roving cursor complicates coalescing and compaction for no clear gain), size-indexed trees (O(log h) but harder to keep deterministic) and the buddy system (power-of-two sizes cause internal fragmentation).

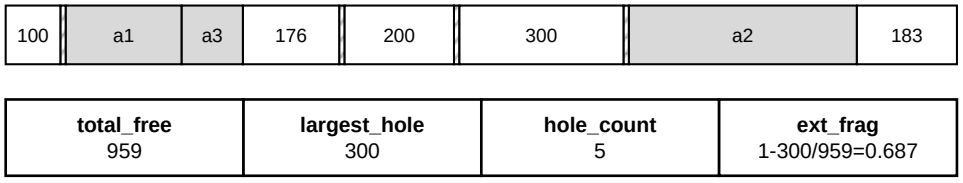
VIII. Fragmentation and Compaction

The problem. After many allocations and frees, free memory is broken into small gaps. A new request may fail even though the *total* free memory is enough, simply because no single gap is big enough. MemForge measures how bad this is, and can fix it by *compaction*: sliding all used pieces to one end so the gaps merge into one big hole, at the cost of copying data.

TABLE IV
Fragmentation metrics

Metric	Definition
total_free, largest_hole, hole_count	sum of holes; largest; number of holes
ext_frag	1 − largest_hole / total_free (0 if none free), in ppm
failure cause	FRAGMENTATION if total_free ≥ s > largest; else CAPACITY
blocks_moved, units_moved	cost of compaction

Ratios are integer parts-per-million, $(p/w)\cdot 10^6 + ((p \bmod w)\cdot 10^6)/w$, so every machine prints the same value (Fig. 15).



request 426 <= 959 free but > 300 largest => cause FRAGMENTATION

Fig. 15. Fragmentation metrics when First-Fit fails the 426-unit request.

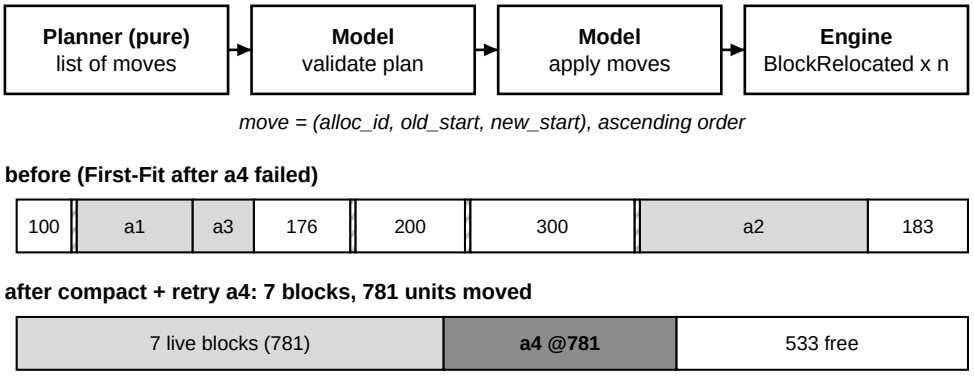


Fig. 16. Compaction pipeline and its effect on AC-2.

Tested guarantees: the live set and total free memory are unchanged, at most one hole remains, block order is preserved, and 5,000 random layouts agree with an independent oracle.

IX. Memory Allocation Examples and Results

TABLE V
AC-1: start address per request (memory 1740)

Req.	Size	FF	BF	WF
a1	212	110	830	1140
a2	417	1140	110	110
a3	112	322	620	1352
a4	426	fail	1140	fail

TABLE VI
AC-1 summary metrics

Metric	FF	BF	WF
holes scanned	23	29	29
mean scanned	1.769	2.231	2.231
used / free	781 / 959	1207 / 533	781 / 959

First- and Worst-Fit fail a4 with 959 units free because the largest hole is 300 (external fragmentation). In AC-2, compaction moves 7 blocks (781 units) into one 959-unit hole [781, 1740); the retried a4 lands at 781, leaving [1207, 1740) free (Fig. 16).

Experiment E1 runs 30 seeded workloads (60/75/90% target utilisation × 10 seeds; 3,000 operations on 1,024 units; sizes 16–128; allocation probability 0.6). Figs. 17–18 plot mean values; Section XV gives 95% confidence intervals.

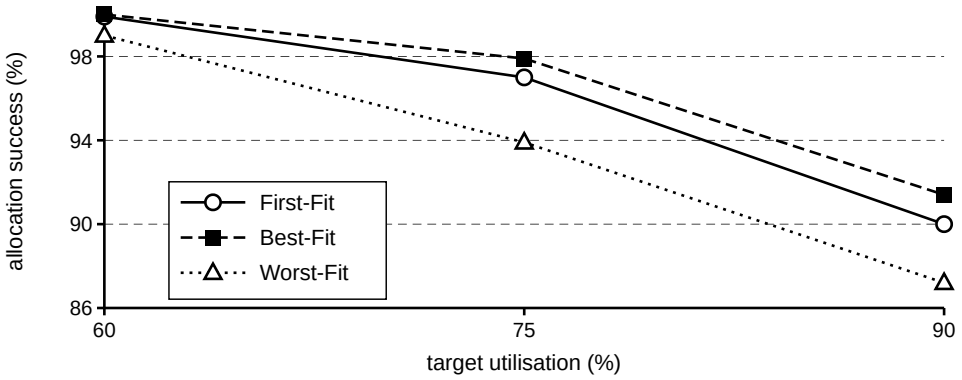


Fig. 17. E1: allocation success rate vs. target utilisation.

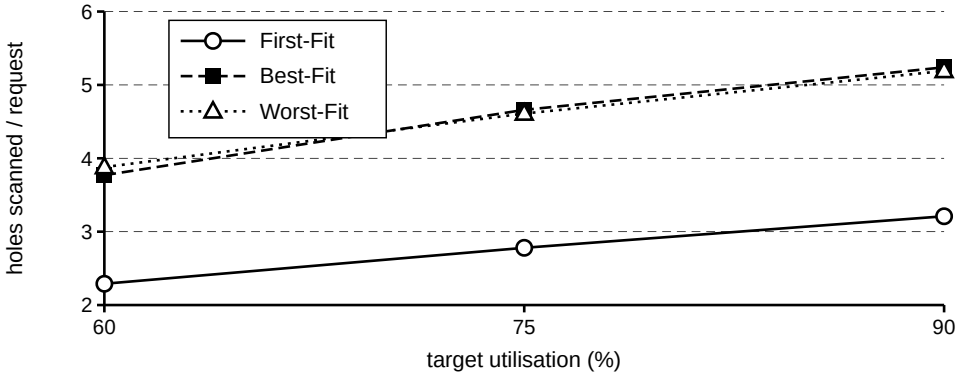


Fig. 18. E1: search cost (mean holes scanned per request).

Best-Fit wins on success and fragmentation, Worst-Fit loses on both, and First-Fit is the cheapest to run. All policies degrade as utilisation rises.

X. Paging Model and Address Translation

Each program sees its own private *virtual* addresses starting at 0. Its memory is cut into pages (here 4096 bytes); RAM is cut into frames of the same size. A per-program *page table* says which frame holds each page. To read an address, the hardware splits it into "which page" and "how far into the page", looks up the frame, and keeps the offset unchanged.

The paging subsystem has four parts: MMU, per-process page tables, an inverted frame table and the paging engine (Fig. 19).

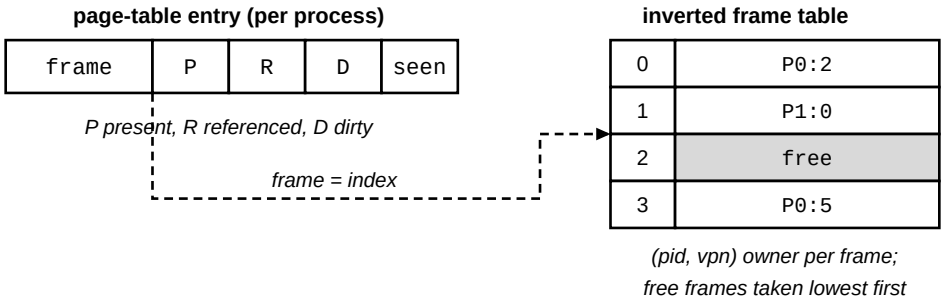


Fig. 19. Page-table entry and inverted frame table.

For a power-of-two page size S, $k = \log_2 S$; vpn = va >> k, offset = va & (S−1) and pa = (f << k) | offset (Fig. 20). The full access path, including hit, fault, eviction and write-back, is Fig. 9. Faults are *compulsory* on a page's first use and *capacity* afterwards. A hit costs O(1), a fault O(F); space is O(P + F).

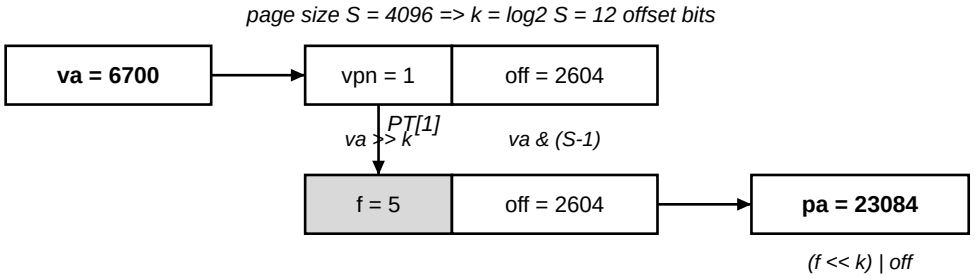


Fig. 20. Translating va = 6700 with page size 4096.

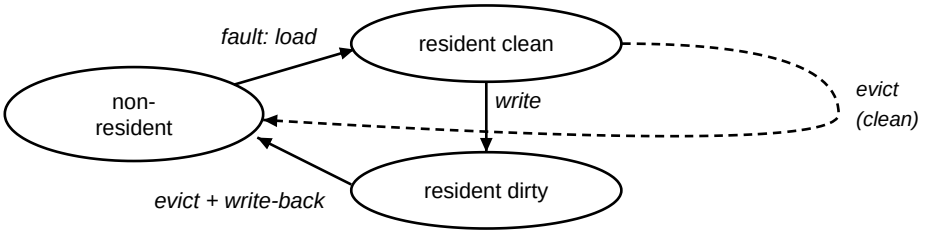


Fig. 21. Page life cycle; dirty victims are written back.

TABLE VII
PAGING INVARIANTS (CHECKED AFTER EVERY STEP)

ID	Invariant
V1	page tables and frame owners agree
V2	every frame is free or owned
V3	non-resident pages have no frame, R or D
V4	resident pages ≤ frames
V5	refs = hits + faults; write-backs ≤ evictions

TABLE VIII
NEGATIVE TESTS (STATE UNCHANGED)

Access	Condition	Result
P0, page 9	outside address space	OUT_OF_RANGE
P9, page 0	no such process	NO_SUCH_PROCESS

XI. FIFO, LRU and Optimal Replacement

When a needed page is not in RAM and every frame is already full, some page must be thrown out to make room. A good choice throws out a page that will not be needed soon; a bad choice throws out a page that is needed again immediately, causing another fault. FIFO, LRU and OPT are three rules for making this choice.

A replacement policy keeps only its own bookkeeping and answers one question: which frame to evict. The engine calls its hooks in a fixed order (Fig. 22); ties go to the lowest frame.

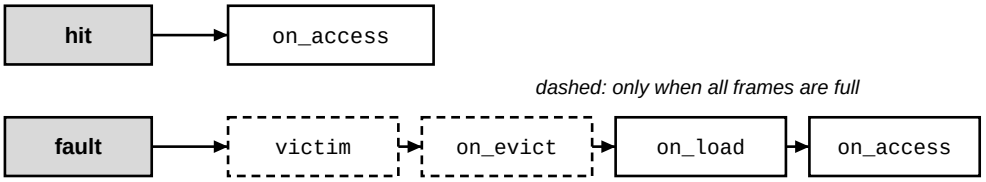


Fig. 22. Order of replacement-policy hooks on a hit and on a fault.

TABLE IX
REPLACEMENT POLICIES COMPARED

Policy	State	Per fault	Future?	Stack?
FIFO	load time / frame	O(F)	no	no
LRU	last use / frame	O(F)	no	yes
OPT	next-use table	O(F)	yes	yes

OPT evicts the page whose next use is furthest away, which can never cause an earlier fault than any other choice [3]; it is therefore the lower bound used to judge FIFO and LRU. Because it needs the whole reference string, the session refuses OPT in live interactive runs.

XII. Page Replacement Worked Examples

The grids below are what `memsim refs --table` prints for the 20-reference string of Silberschatz with 3 frames (AC-3). Shaded cells mark the frame loaded on a fault.

ref	**7**	**0**	**1**	**2**	**0**	**3**	**0**	**4**	**2**	**3**	**0**	**3**	**2**	**1**	**2**	**0**	**1**	**7**	**0**	**1**
f0	7	7	7	2	2	2	2	4	4	4	0	0	0	0	0	0	0	7	7	7
f1		0	0	0	0	3	3	3	2	2	2	2	2	1	1	1	1	1	0	0
f2			1	1	1	1	0	0	0	3	3	3	3	3	2	2	2	2	2	1
fault	**F**	**F**	**F**	**F**		**F**	**F**	**F**	**F**	**F**	**F**			**F**	**F**			**F**	**F**	**F**

Fig. 23. FIFO: 15 faults (6 compulsory, 9 capacity).

ref	7	0	1	2	0	3	0	4	2	3	0	3	2	1	2	0	1	7	0	1
f0	7	7	7	2	2	2	2	4	4	4	0	0	0	1	1	1	1	1	1	1
f1		0	0	0	0	0	0	0	0	3	3	3	3	3	3	0	0	0	0	0
f2			1	1	1	3	3	3	2	2	2	2	2	2	2	2	2	7	7	7
fault	F	F	F	F		F		F	F	F	F			F		F		F		

Fig. 24. LRU: 12 faults (6 compulsory, 6 capacity).

ref	7	0	1	2	0	3	0	4	2	3	0	3	2	1	2	0	1	7	0	1
f0	7	7	7	2	2	2	2	2	2	2	2	2	2	2	2	2	2	7	7	7
f1		0	0	0	0	0	0	4	4	4	0	0	0	0	0	0	0	0	0	0
f2			1	1	1	3	3	3	3	3	3	3	3	1	1	1	1	1	1	1
fault	F	F	F	F		F		F			F			F				F		

Fig. 25. OPT: 9 faults (6 compulsory, 3 capacity).

All three evict page 7 at reference 4. They first diverge at reference 6 (page 3): FIFO evicts page 0, which is needed immediately, while LRU and OPT evict page 1.

A. Belady's Anomaly (AC-4)

ref	1	2	3	4	1	2	5	1	2	3	4	5
f0	1	1	1	4	4	4	5	5	5	5	5	5
f1		2	2	2	1	1	1	1	1	3	3	3
f2			3	3	3	2	2	2	2	2	4	4
fault	F	F	F	F	F	F	F			F	F	

Fig. 26. FIFO, 3 frames: 9 faults.

ref	1	2	3	4	1	2	5	1	2	3	4	5
f0	1	1	1	1	1	1	5	5	5	5	4	4
f1		2	2	2	2	2	2	1	1	1	1	5
f2			3	3	3	3	3	3	2	2	2	2
f3				4	4	4	4	4	4	3	3	3
fault	F	F	F	F			F	F	F	F	F	F

Fig. 27. FIFO, 4 frames: 10 faults.

With 4 frames, from the seventh reference on FIFO always evicts exactly the page needed next, so every reference faults.

XIII. Paging Results and Comparison

TABLE X
E4 configuration

Parameter	Value
virtual pages / page size	32 / 4096 B
references per run	2,000
working set / phase length	6 pages / 200 refs
out-of-set probability	5%
seeds / frames	10 / 1–16

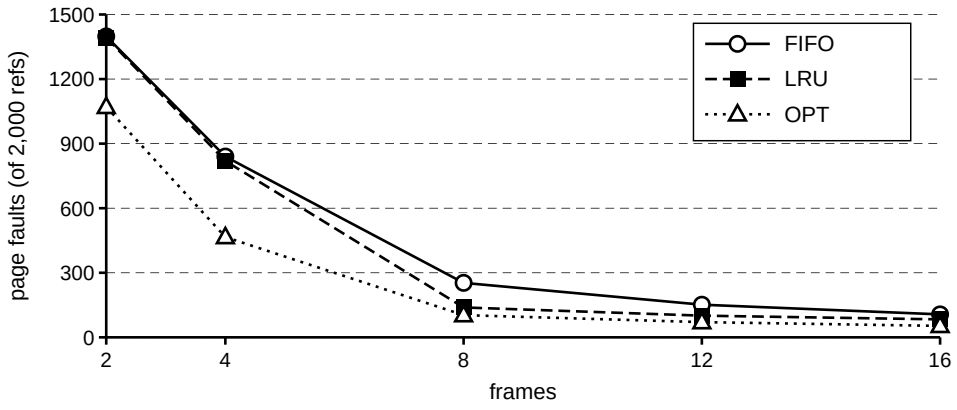


Fig. 28. E4: mean page faults vs. frames (10 seeds).

TABLE XI
E4 PAGE FAULTS, MEAN ± 95% CI

Frames	FIFO	LRU	OPT
2	1398.2 ± 21.2	1393.0 ± 18.8	1070.8 ± 13.9
4	840.0 ± 15.9	819.6 ± 19.9	466.0 ± 9.3
8	253.1 ± 16.5	138.9 ± 10.1	103.2 ± 5.9
12	151.7 ± 9.7	100.7 ± 6.0	70.7 ± 4.0
16	106.7 ± 9.0	83.5 ± 5.5	53.2 ± 3.2

Smooth locality hides non-monotonic behaviour, but the AC-4 sweep exposes it: FIFO gives 12, 12, 9, 10, 5, 5 faults for 1–6 frames while LRU and OPT never rise (Fig. 5). In practice OPT is unimplementable, FIFO is simple but anomalous, and LRU is strong under locality but costly, which is why real systems approximate it with Clock.

XIV. Workloads and Performance Metrics

A *workload* is simply a script of memory requests (allocate, free, access a page) that is fed to the simulator. A *metric* is a number measured while it runs, such as how many requests succeeded or how many page faults occurred. Using the same workload for every policy makes the comparison fair.

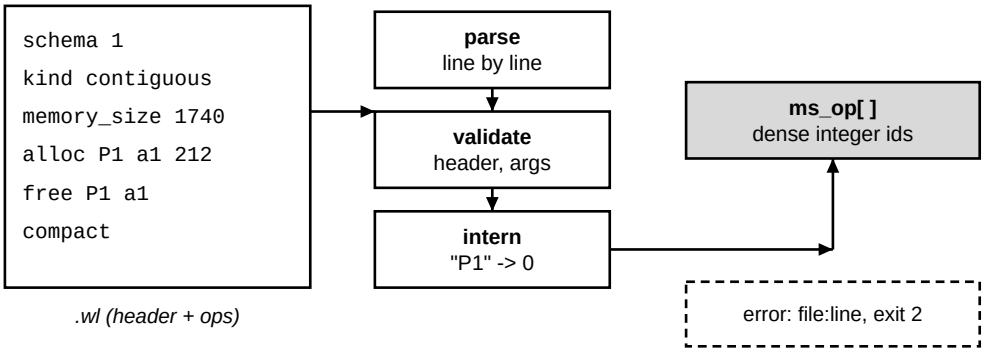


Fig. 29. Workload file format and loader pipeline.

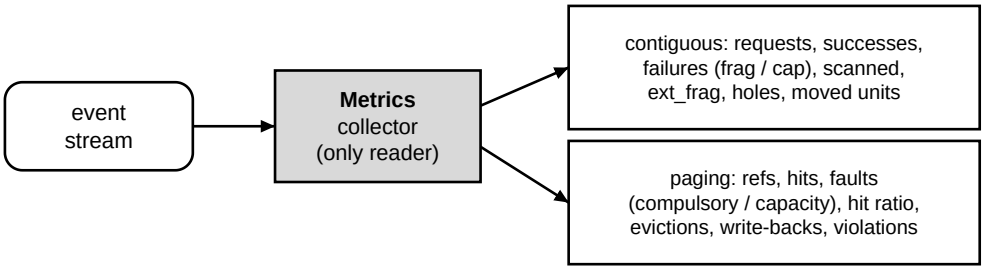


Fig. 30. One collector derives every metric from the event stream.

XV. Benchmarking and Performance Analysis



E1: 3 utilisations x 10 seeds = 30 workloads, 3,000 ops, 1,024 units, sizes 16-128

E4: 10 seeds x 2,000 refs, 32 pages, working set 6, phase every 200, 5% outside

paired design: every policy replays the same file per seed

scripts/experiments/run_all.sh regenerates everything

Fig. 31. Experiment pipeline (paired, seeded, reproducible).

TABLE XII
E1 RESULTS WITHOUT COMPACTION, MEAN ± 95% CI

Util.	Policy	Success %	Frag. fail.	Scanned	Ext. frag.
60%	FF	99.9 ± 0.0	1.2 ± 0.7	2.29	0.359
60%	BF	100.0 ± 0.0	0.3 ± 0.3	3.77	0.322
60%	WF	99.0 ± 0.2	14.3 ± 3.1	3.88	0.546
75%	FF	97.0 ± 0.3	44.6 ± 4.7	2.78	0.460
75%	BF	97.9 ± 0.2	32.0 ± 2.8	4.66	0.419
75%	WF	93.9 ± 0.5	92.4 ± 6.8	4.61	0.597
90%	FF	90.0 ± 0.6	148.6 ± 8.6	3.21	0.516
90%	BF	91.4 ± 0.4	127.7 ± 7.0	5.24	0.473
90%	WF	87.2 ± 0.7	192.0 ± 11.1	5.19	0.629

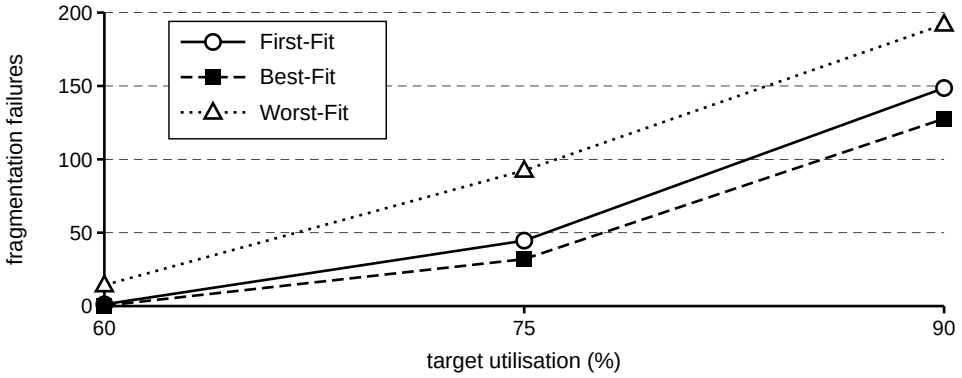


Fig. 32. E1: requests failed by fragmentation despite enough free memory.

Quality: Best-Fit has the highest success and lowest fragmentation at every level. **Cost:** First-Fit scans 2.3–3.2 holes vs. 3.8–5.2. **Pressure:** at 90%, fragmentation failures reach 149 (FF), 128 (BF) and 192 (WF). **Reproducibility:** every cell regenerates from the generator, the seed and the git revision.

XVI. TESTING STRATEGY AND CORRECTNESS

Testing answers one question: does the simulator give the right answers? We check it against textbook examples whose answers are known, check internal consistency rules after every single step, run it twice to confirm identical output, and run memory-error detectors over the C code.

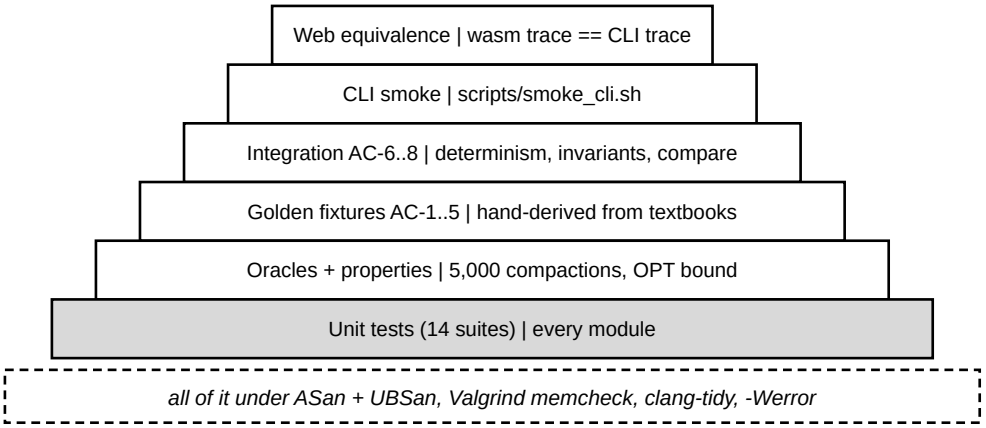


Fig. 33. Independent test layers; a bug missed by one is caught by another.

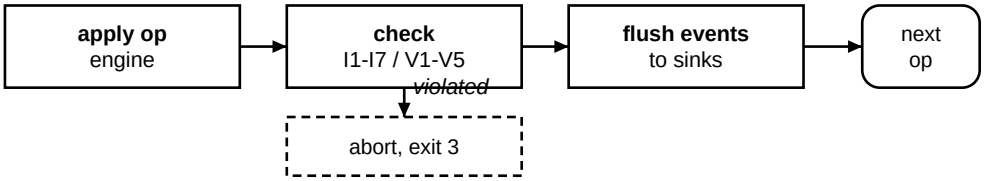


Fig. 34. Checked mode: invariants after every operation.

Randomised oracles: 5,000 compactions match a unit-by-unit oracle; no policy beats OPT; LRU and OPT never show Belady's anomaly; every fixture run twice gives equal SHA-256 trace digests (AC-6). Expected values are written into each test before it is compared with the engine, taken from the textbook [1] or from Belady's paper [4].

XVII. Test Cases and Test Results

TABLE XIII
Acceptance criteria (verify_report.sh)

ID	Scenario	Expected	Result
AC-1	five-hole placement	FF/WF fail a4, BF places all	PASS
AC-2	compaction + retry	7 blocks, 781 units; [1207,1740)	PASS
AC-3	20 refs, 3 frames	FIFO 15, LRU 12, OPT 9	PASS
AC-4	Belady sweep	FIFO 12,12,9,10,5,5	PASS
AC-5	negative tests	3 rejections, 2 violations	PASS
AC-6	determinism	6/6 traces identical	PASS
AC-7	invariants	17/17 suites	PASS
AC-8	compare table	shows fail FRAGMENTATION	PASS

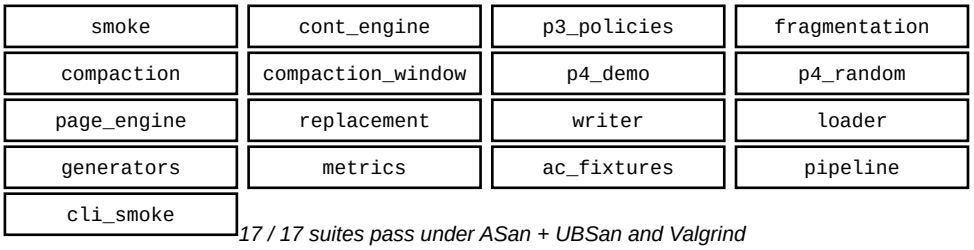


Fig. 35. Automated test suites (all passing).

XVIII. User Interface and Workflow

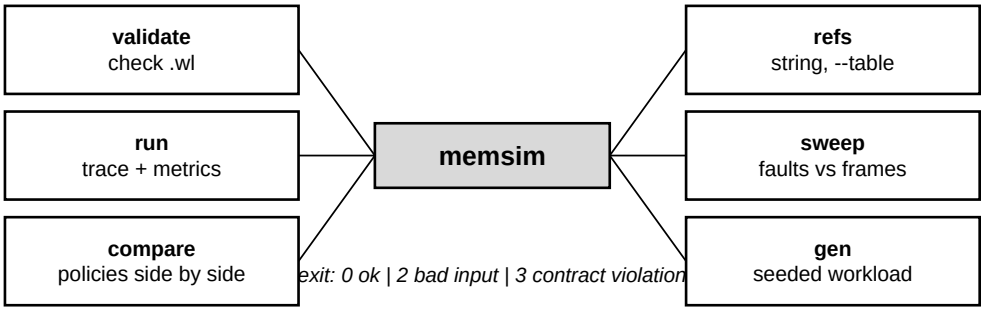


Fig. 36. The memsim command-line interface.

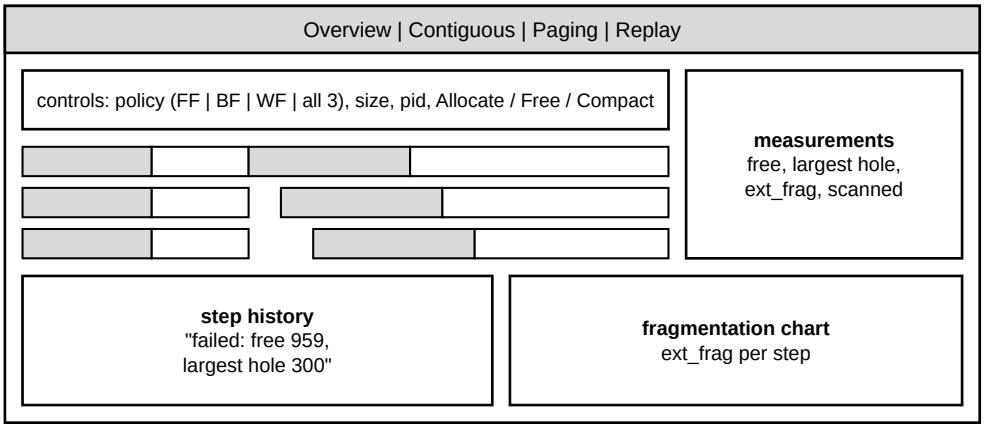


Fig. 37. Layout of the Contiguous page of the web laboratory.

The lab has four pages: Overview (live run and six classic scenarios), Contiguous, Paging (frames grid, page tables, translation, Belady curve) and Replay (step through any `memsim run --trace` file). Neither interface contains allocation or replacement logic.

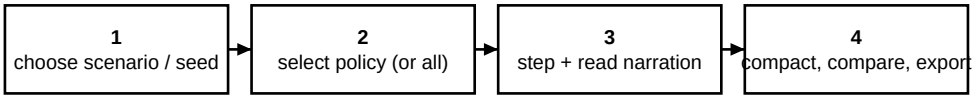


Fig. 38. Typical workflow in the lab.

XIX. Execution Examples and Output

Unedited terminal output (commit ab5d17d).

```
$ memsim compare ac1_silberschatz.wl \
    --policies first-fit,best-fit,worst-fit
id  size  first-fit           best-fit  worst-fit
a1  212   110                 830       1140
a2  417   1140                110       110
a3  112   322                 620       1352
a4  426   fail FRAGMENTATION  1140      fail FRAGMENTATION
```

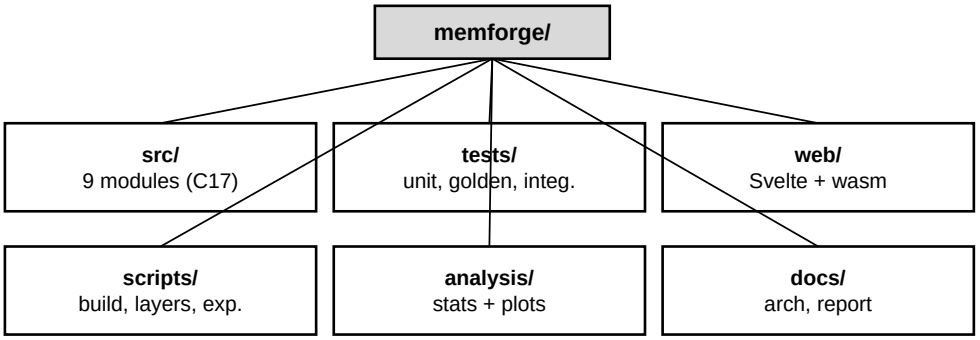
```
$ memsim run ac2_compact.wl --policy first-fit
compactions 1   blocks moved 7   units moved 781
ext frag final 0.000000   hole count final 1
total free final 533   largest hole final 533
```

```
$ memsim sweep ac4_belady.wl --policy fifo --frames 1..6
1 12
2 12
3 9
4 10 *
5 5
6 5
```

```
$ memsim run ac5_contiguous.wl  -> rejections 3
$ memsim run ac5_paging.wl      -> violations 2
access P0 @6700 -> AddressTranslated vpn 1,
    offset 2604, frame 1, pa 6700
run summary: refs 4 hits 1 faults 3 evictions 1
```

```
$ memsim gen alloc-uniform --seed 7 -o w.wl
schema 1 / kind contiguous / memory_size 1024 /
gen_name alloc-uniform / gen_seed 7 / gen_ops 2000
```

XX. Repository Structure and Development Conventions



~6,900 lines C in src/, ~4,000 lines C tests, ~3,100 lines web

Fig. 39. Top-level repository layout.

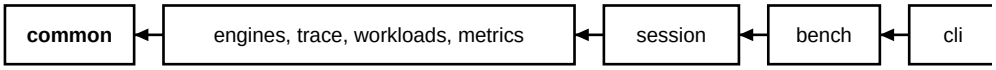


Fig. 40. Allowed include direction between module layers.

TABLE XIV
Conventions and ownership

Rule	Practice
naming	prefix ms_ / MS_, e.g. ms_cont_engine_alloc
format	clang-format, .editorconfig
ownership	CODEOWNERS, roles P1–P10
interfaces	change via RFC, lead approves
commits	pre-commit: format + layering

XXI. Build System and Continuous Integration

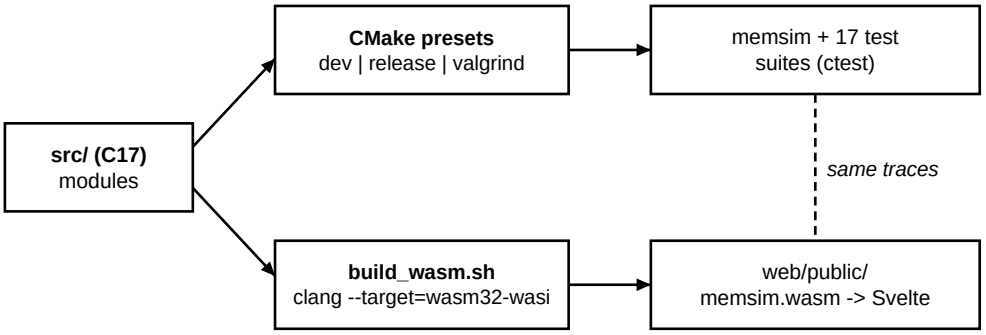


Fig. 41. Build orchestration: native presets and the WebAssembly build.

TABLE XV
CMAKE PRESETS

Preset	Configuration	Use
dev	Debug + ASan + UBSan	daily work, main tests
release	-O2	experiments, binary
valgrind	Debug	memcheck on every test

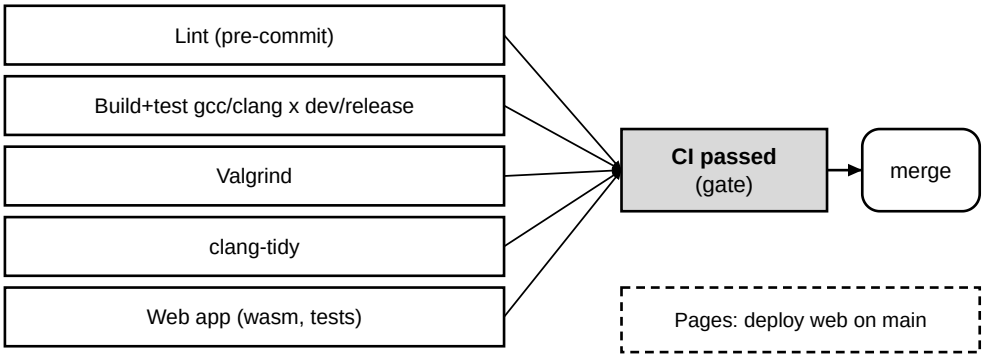
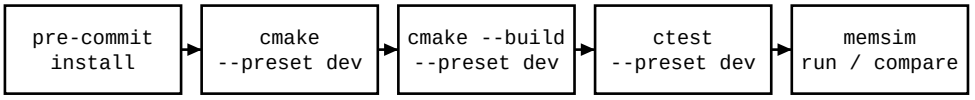


Fig. 42. GitHub Actions jobs and the single merge gate.

All code builds as C17 with `-Wall -Wextra -Wpedantic -Wshadow -Wconversion -Wsign-conversion`, and with `-Werror` in CI.

XXII. Setup Instructions and Documentation

Requirements: Linux or macOS (Windows via WSL), GCC or Clang, CMake ≥ 3.25, Ninja, pre-commit; Python 3 + matplotlib for experiments; Node.js and Clang/lld/wasi-libc for the web app.



web: build_wasm.sh -> npm ci -> npm run dev | npm test

Fig. 43. Build, test and run sequence.

```
$ cmake --preset release && cmake --build --preset release
$ bash scripts/experiments/run_all.sh
$ bash scripts/report/verify_report.sh
```

TABLE XVI
DOCUMENTATION

Document	Content
README.md	overview, quick start, commands
docs/ARCHITECTURE.md	requirements, design, ADRs
docs/report/	report, tables, figures
analysis/README.md	statistics and plots
.github/	PR, bug and RFC templates

A hosted copy of the lab is at memforge.aryagupta-dev.xyz.

XXIII. Conclusion and Future Scope

MemForge meets its goals: a deterministic, modular C17 simulator that reproduces the textbook examples exactly and compares policies fairly. Best-Fit gives the best success and lowest fragmentation, First-Fit is cheapest, Worst-Fit is worst on both; OPT < LRU < FIFO at every frame count, and only FIFO shows Belady's anomaly.

TABLE XVII
GOALS AND EVIDENCE

Area	Delivered	Evidence
Contiguous	FF/BF/WF, split, coalesce, compaction	AC-1, AC-2, E1
Paging	translation, FIFO/LRU/OPT, sweep	AC-3–5, E4
Determinism	byte-identical traces	AC-6
Correctness	invariants, oracles	AC-7, 17/17
Delivery	CLI + wasm from one code base	web equivalence test

Lessons: determinism must be designed in (integer state, logical clocks, tie-breaks); separating decisions from mutation made policies trivially testable; compiling the real engine to WebAssembly removed demo drift; automated layering checks caught boundary violations early.

Future work: Next-Fit, buddy allocation, alignment; Clock, LFU, a TLB and two-level tables; local replacement; Knuth's fifty-percent rule; real program traces; fuzzing and a coverage gate.

References

[1] A. Silberschatz, P. B. Galvin and G. Gagne, *Operating System Concepts*, 10th ed. Hoboken, NJ, USA: Wiley, 2018.

[2] A. S. Tanenbaum and H. Bos, *Modern Operating Systems*, 4th ed. Boston, MA, USA: Pearson, 2015.

[3] L. A. Belady, “A study of replacement algorithms for a virtual-storage computer,” *IBM Syst. J.*, vol. 5, no. 2, pp. 78–101, 1966.

[4] L. A. Belady, R. A. Nelson and G. S. Shedler, “An anomaly in space-time characteristics of certain programs running in a paging machine,” *Commun. ACM*, vol. 12, no. 6, pp. 349–353, Jun. 1969.

[5] D. E. Knuth, *The Art of Computer Programming, Vol. 1: Fundamental Algorithms*, 3rd ed. Reading, MA, USA: Addison-Wesley, 1997, sec. 2.5.

[6] P. J. Denning, “The working set model for program behavior,” *Commun. ACM*, vol. 11, no. 5, pp. 323–333, May 1968.

[7] R. L. Mattson, J. Gecsei, D. R. Slutz and I. L. Traiger, “Evaluation techniques for storage hierarchies,” *IBM Syst. J.*, vol. 9, no. 2, pp. 78–117, 1970.

[8] D. Blackman and S. Vigna, “Scrambled linear pseudorandom number generators,” *ACM Trans. Math. Softw.*, vol. 47, no. 4, art. 36, 2021.

[9] *Information Technology—Programming Languages—C*, ISO/IEC 9899:2018.

10] WebAssembly Community Group, *WebAssembly Core Specification*, W3C.

11] Kitware, *CMake Reference Documentation*. [Online]. Available: https://cmake.org/documentation/

12] *Svelte Documentation*, https://svelte.dev/docs; *Valgrind User Manual*, https://valgrind.org/docs/manual/.

Appendix

TABLE XVIII
Complete AC-1 comparison metrics

Metric	FF	BF	WF
requests / successes	13 / 12	13 / 13	13 / 12
failures (frag.)	1	0	1
frees	5	5	5
scanned total / mean	23 / 1.769	29 / 2.231	29 / 2.231
ext_frag final / mean	0.687 / 0.246	0.674 / 0.212	0.687 / 0.245
holes / free / largest	5 / 959 / 300	5 / 533 / 174	5 / 959 / 300
used / live blocks	781 / 7	1207 / 8	781 / 7

TABLE XIX
Paging metrics for the textbook strings

Run	Hits	Faults	Comp.	Cap.	Hit ratio	Evict.
AC-3 FIFO	5	15	6	9	0.250	12
AC-3 LRU	8	12	6	6	0.400	9
AC-3 OPT	11	9	6	3	0.550	6
AC-4 FIFO	3	9	5	4	0.250	6
AC-4 LRU	2	10	5	5	0.167	7
AC-4 OPT	5	7	5	2	0.417	4

```
{"seq":40,"step":18,"type":"PlacementDecided",
 "data":{"policy":"first-fit","requested":426,
  "hole":null,"scanned":5}}
{"seq":41,"step":18,"type":"AllocationFailed",
 "data":{"pid":"P4","id":"a4","requested":426,
  "cause":"FRAGMENTATION","total_free":959,
  "largest_hole":300}}
```

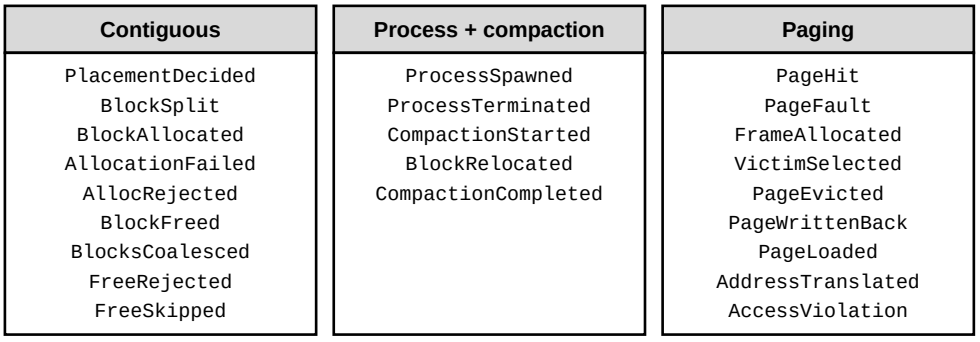


Fig. 44. Event catalogue (23 types).